Search-Based Genetic Optimization for Deployment and Reconfiguration of Software in the Cloud

Sören Frey, Florian Fittkau, and Wilhelm Hasselbring
Software Engineering Group
Kiel University
24118 Kiel, Germany
{sfr, ffi, wha}@informatik.uni-kiel.de

Abstract—**Migrating existing enterprise software to cloud platforms involves the comparison of competing cloud deployment options (CDOs). A CDO comprises a combination of a specific cloud environment, deployment architecture, and runtime reconfiguration rules for dynamic resource scaling. Our simulator CDOSim can evaluate CDOs, e.g., regarding response times and costs. However, the design space to be searched for well-suited solutions is extremely huge. In this paper, we approach this optimization problem with the novel genetic algorithm CDOXplorer. It uses techniques of the search-based software engineering field and CDOSim to assess the fitness of CDOs. An experimental evaluation that employs, among others, the cloud environments Amazon EC2 and Microsoft Windows Azure, shows that CDOXplorer can find solutions that surpass those of other state-of-the-art techniques by up to 60%. Our experiment code and data and an implementation of CDOXplorer are available as open source software.**

Index Terms—**Cloud computing, Search-based software engineering, Deployment optimization**

I. INTRODUCTION

The disruptive cloud computing paradigm paves the way for approaching the long-desired idea of utility computing [1]. In the last years, it gained considerable attention from industry and in academia. Along with the steadily increasing interest in cloud computing there also emerged an enormous demand for leveraging cloud technologies for existing systems [2]. However, migrating and deploying enterprise software to the cloud still entails a wealth of challenges and potential pitfalls. For example, it is tedious to select an adequate cloud environment and the best-suited virtual machine (VM) instance types—with regard to inevitable trade-offs between costs and performance—from the plethora of available cloud offerings. Then, the application and deployment architecture have to be reworked to conform with the chosen cloud and to enable compliance with defined service level agreements (SLAs) and the included quality of service (QoS) stipulations. To exploit the cloud's elasticity and the usually employed pay-per-use model, it is necessary to implement and calibrate reconfiguration rules for cost-efficient dynamic resource scaling according to observed usage patterns. All of those design decisions form a multitude of cloud deployment options (CDOs) that need to be explored for well-suited candidates.

Unfortunately, techniques for automatically evaluating all CDOs do not exist and a comprehensive manual analysis is most often inapt due to time and budget constraints [3].

Furthermore, as we also intend to support distributed systems, we consider the QoS-aware composition of software components that run on one node as a single service that is provided to an arbitrary number of components on other nodes. Such deployment optimization problems are intractable as they are known to be NP-hard [4]. In our previous work, we introduced the simulation tool CDOSim [5] that implements a phase of our cloud migration approach CloudMIG [6, 7]. CDOSim facilitates the simulation of CDOs for determining their respective response times, costs, and SLA violations.

We integrated CDOSim into our tool CloudMIG Xpress that provides support for CloudMIG. With CloudMIG Xpress, CDOs can be manually configured and simulated on the basis of a reverse-engineered architectural system model with monitored or synthetic workload. However, the design space that spans for all possible CDOs is huge, the elements of a single CDO exhibit complex non-linear interdependencies, and CDO simulation runs are very time-consuming and can take from a few minutes to several hours. Hence, simulating a great number of CDOs is most often still not a viable option and it is therefore very likely that a suboptimal solution is chosen. Moreover, there usually exists no single CDO that causes the lowest costs along with the lowest average response times and the lowest number of SLA violations. Thus, a potential cloud user is interested in automatically finding the most adequate trade-off solutions among which the CDO candidate can be selected that best suites the user's specific needs. The set of these most adequate trade-off solutions constitutes a pareto optimum. The included CDOs cannot be improved concerning one objective without deteriorating another objective, e.g., considering a trade-off between costs and response times.

In this paper, we present the genetic algorithm CDOXplorer that explores the CDO search-space on the basis of automatically extracted architectural models and approximates the corresponding pareto optimum. Similar problems are addressed by methods of the search-based software engineering field, where genetic algorithms are widely used [8]. CDOXplorer is implemented in our tool CloudMIG Xpress and supports IaaS-based cloud environments [9], where the most often used building blocks are VMs. In general, genetic algorithms group the candidates—so-called individuals—in populations and use a fitness function to assess the candidates. Then, the best-suited individuals are selected. They reproduce through so-

called mutation and crossover operations and after several generations, the individuals that inherited superior properties become dominant. To assess the fitness of CDOs, CDOXplorer uses simulation runs of CDOSim to restrict the search-space and to steer the exploration towards promising CDOs. Thus, CDOSim is no longer used only for analyses, but for design purposes as well. By incorporating CDOSim, CDOXplorer is a member of the simulation-based optimization class [10].

Here, the evaluation of the used fitness function is, in contrast to most genetic algorithms, very expensive and requires strict limitations regarding the population size and number of included generations. CDOXplorer not only optimizes the allocation of software components to VMs, but also searches for reconfiguration rules that are aligned with the cloud's elasticity and the specific performance and pricing models of the available cloud environments. A common challenge in the design of genetic algorithms becomes apparent as they do not guarantee to converge to a global optimum, especially, if a low number of generations is used. Hence, we experimentally evaluate the applicability and convergence properties in comparison to other well-known search and optimization algorithms. We report on case studies that employ an open source enterprise resource planning (ERP) system to be deployed on the cloud environments Amazon EC2, an Eucalyptus cluster as private cloud, and Microsoft Windows Azure. In summary, our main contributions are:

- A simulation-based genetic algorithm (CDOXplorer) for finding near-optimal cloud deployment architectures and runtime reconfiguration rules for enterprise software
- An implementation of CDOXplorer within the scope of our open source tool CloudMIG Xpress, that utilizes models which can almost all be extracted automatically
- Extensive experiments that employ well-known clouds show that CDOXplorer can find results that surpass those of other state-of-the-art techniques by up to 60%

CloudMIG Xpress and our experiment code and data are available online as open source software such that interested researchers may repeat or extend our experiments.[1]

The paper is organized as follows. We provide a motivating example in Section II. The required input models for our genetic algorithm are described in Section III and the output models it produces are explained in Section IV. Section V details CDOXplorer. Then, the experimental evaluation of CDOXplorer is described in Section VI. Section VII discusses related work, before our conclusions are drawn in Section VIII.

II. MOTIVATING EXAMPLE

We consider the example deployment of a software system that is shown on the left side of Figure 1. This system should be moved to the cloud environment Amazon EC2. It consists of eight software components that are currently deployed to three interconnected on-premise server machines. Those machines are also called *status-quo nodes* as they constitute elements of the deployment architecture that describes the
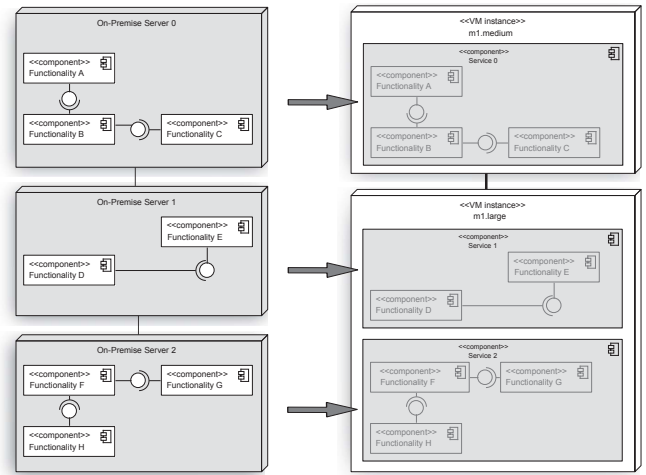
[1]http://www.cloudmig.org

Fig. 1: Mapping on-premise servers and deployed components (left) to atomic services in a basic CDO example (right).

status-quo assignment of components to physical machines. As mentioned earlier, we regard all of the components that are deployed to a single status-quo node as a separate *service*. A service is an atomic unit concerning the allocation to a VM. This design decision was made to (1) prevent a further explosion in the number of combinations and CDOs that have to be searched, and to (2) render pervasive changes unnecessary that may be required when distributing tightly connected components over different VMs.

Thus, the three status-quo nodes and the deployed components shown in the left part of Figure 1 result in the three services *Service 0* to *Service 2* in the right part of Figure 1 that exhibit a similar assignment of components. As can be seen in the lower right part of Figure 1, a single VM can host multiple services. In our example, it was decided to consolidate *Service 1* and *Service 2* into a joint VM that is started with Amazon EC2's VM instance type *m1.large*. Such VM instance types describe the hardware resources that are available to VMs. For example, at the time of writing, Amazon EC2's *m1.large* VM instance type provides 7.5 GB memory and two virtual cores that together provide approximately the CPU capacity of four 1.7 GHz Xeon processors from 2006. The remaining *Service 0* in Figure 1 is deployed to an own VM that builds upon the *m1.medium* VM instance type.

The basic CDO of Figure 1 is now given by the number of chosen VM instances, the assignment of services to these VM instances, and the selection of a VM instance type for each VM instance. Furthermore, up to this point, runtime reconfiguration rules are omitted for the sake of simplicity.

Reasoning about the broader array of all potential CDOs for, e.g., just the single cloud environment Amazon EC2, 10 of its VM instance types, and no reconfiguration rules, reveals the general complexity of CDO analysis. When assuming up to three VM images that contain combinations of the three services and that up to two VMs can be started from a VM image, these restrictive settings already yield 2,744,000 CDO candidates. Without using potent heuristics, all CDOs would have to be simulated for reliably finding competitive solutions.

III. INPUT MODELS

In this section, we briefly describe the four input models that have to be provided to CDOXplorer so it can find well-suited cloud deployment models and reconfiguration rules. We provide tool support for creating these models. The *architectural model*, *status-quo deployment model*, *workload profile*, and *cloud profile* are described in the following.

An *architectural model* can be automatically extracted by CloudMIG Xpress from an existing system's source code. Currently, we support Java, C#, and Python by incorporating, among others, the reverse-engineering tool MoDisco [11]. CloudMIG Xpress generates architectural models that conform to the ISO/IEC 19506 standard Knowledge Discovery Meta-Model [12] (KDM) that was developed by the Object Management Group (OMG). With KDM, several aspects of software systems can be modeled in a language-independent way, for example, the runtime platforms and source code elements.

To describe the current deployment, a *status-quo deployment model* is designed with the integrated editor. Here, KDM code elements can be assigned to status-quo nodes. To measure and specify the performance capabilities of these computing nodes, we introduced a benchmark that measures the mega integer plus instructions per second (MIPIPS) [5]. The computing power must be specified to enable CDOSim to interpret a *workload profile* that describes service calls and response times and can be imported from monitoring log files with historical usage data. Currently, we support Kieker [13] log files, but additional monitoring log formats can easily be incorporated via plug-ins. If no real monitoring data is present, CloudMIG Xpress also allows for the definition of synthetic workload profiles. For specifying cloud environments, so-called *cloud profiles* are used that, e.g., describe a VM instance types' performance capabilities. MIPIPS can also be measured by executing the mentioned benchmark for each VM instance type of a cloud environment. We plan to develop and integrate a public repository such that cloud profiles can be exchanged.

IV. OUTPUT MODELS

This section explains the result models that are included in a CDO as delivered by CDOXplorer. The CDOs specify a cloud deployment model and a set of reconfiguration rule models. In IaaS-based cloud environments, resources can be dynamically acquired and released by executing *reconfiguration actions* to counteract under- and over-provisioning and to ensure the compliance with specified SLAs. Two reconfiguration actions together define one of the scaling types *vertical scaling* or *horizontal scaling*. Considering vertical scaling, the reconfiguration actions *scale-up* and *scale-down* are available. Scaling up adds more resources to a VM, e.g., a further CPU or more memory, whereas scaling down removes resources. Horizontal scaling employs VMs that use the same VM instance type. VM instances are added (*scale-out*) or shut down (*scale-in*).

We now describe the basic structure of CDOs that is shown in Figure 2. A *Cloud Deployment Option* refers to a single *Cloud Environment*, that is derived from a cloud profile, and contains so-called *Node Configurations*. A *Node Configuration*
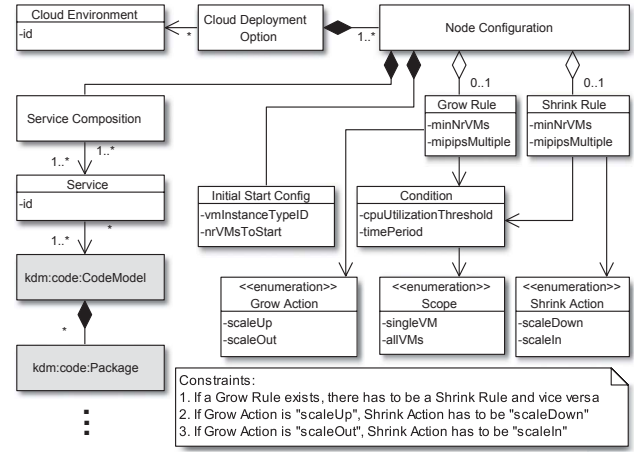

Fig. 2: Basic structure of CDOs.

describes specifics of a VM instance. For example, an included *Service Composition* container refers to the *Service*s that are deployed on this VM. To link *Service*s with the architectural input model, they reference parts of the previously extracted KDM elements. Furthermore, an *Initial Start Config* specifies the VM instance type that should be used for a VM and also the number of VM instances that have to be started initially with this configuration. Moreover, a *Node Configuration* may contain a *Grow Rule* together with a *Shrink Rule*. They represent basic parts of a reconfiguration rule and, from a high level view, determine how and when computing power is added (*Grow Rule*) or removed (*Shrink Rule*). Theses rules also specify a minimum number of VMs that have to be present and refer to a *Grow Action* or *Shrink Action* that define the reconfiguration actions that are used for scaling. A combination of two actions has to comply with the previously explained scaling types.

When applying vertical scaling, the *mipipsMultiple* attribute of *Grow Rule*s and *Shrink Rule*s becomes relevant. It implicitly specifies the VM instance type of the newly started VM, which is given by multiplying the MIPIPS value of the current VM with *mipipsMultiple* and rounding to the nearest MIPIPS value of any (the intended) VM instance type. A *Condition* constitutes a trigger for executing the reconfiguration action with the help of a CPU utilization threshold and a time period. For example, concerning a scale-out action, an additional VM could be started when the CPU utilization is above 80% for at least 20 minutes. The *Scope* element would then define if the 80% refers to the specific VM or to the average of all VMs that were started from the corresponding *Node Configuration*.

V. OUR SIMULATION-BASED GENETIC ALGORITHM

This section details our simulation-based genetic algorithm CDOXplorer. First, Section V-A describes its basic design. Then, Section V-B formalizes the optimization problem that is tackled by CDOXplorer. Sections V-C and V-D detail the crossover and mutation operator of CDOXplorer, respectively.

A. Basic Design

The central purpose of the genetic algorithm CDOXplorer is to efficiently find well-suited CDOs concerning a set of

Fig. 3: Compound chromosome overview. Gray boxes: chromosomes, white boxes: genes (listed in Table I). `1`, `+`, `?` in the boxes' upper left corner indicate that the elements occur exactly once, at least once, and at most once, respectively.

CDO1	7	0	3	2	0	4	3	1	0.7	10	0	2	0.5	1	0.4	15	1	2	1	4
	CE	SE	IT	NI	GA	G1	G2	G3	G4	G5	SA	S1	S2	S3	S4	S5	SE	IT	NI	

CDO2	7	0	8	1	1	2	9	1	0	1	2.3	0	0.8	20	0	1	0.5	1	0.4	45
	CE	SE	IT	NI	SE	SE	IT	NI	GA	G1	G2	G3	G4	G5	SA	S1	S2	S3	S4	S5

CDO3	8	0	1	2	14	2	0	2	1.4	0	0.85	5	0	3	0.9	1	0.6	35
	CE	SE	SE	SE	IT	NI	GA	G1	G2	G3	G4	G5	SA	S1	S2	S3	S4	S5

Fig. 4: CDO examples encoded as genotypes.

arbitrary IaaS-based cloud environments. In general, genetic algorithms mimic evolutionary processes that describe the advancement of populations over several generations. Evolutionary concepts, such as the survival of the fittest and inheritance of properties that turn out to be advantageous, are included and utilized as optimization techniques. Genetic algorithms are usually used in the context of multi-objective optimization problems. With CDOXplorer, we consider the objectives response times, costs, and SLA violations of a CDO.

As there usually exists no single global optimum, genetic algorithms aim to iteratively approximate the *pareto optimal set* (also referred to as *pareto optimum*). This pareto optimal set is a subset of all individuals that includes all *pareto optimal* individuals, i.e., individuals for which the improvement of one objective (e.g., lower costs) would inevitably lead to a deterioration of another objective (e.g., higher response times). As a consequence, the individuals contained in a pareto optimum constitute trade-off solutions that have to be inspected manually. In general, individuals are compared with so-called fitness functions. After many generations, the fittest individuals become dominant. The reproduction of each generation includes the following four basic steps S1-S4 [8]:

- S1 Select parents
- S2 Recombine parents (crossover)
- S3 Mutate offspring
- S4 Evaluate offspring's fitness

The first step S1 selects individuals for reproduction. For CDOXplorer, S1 is based on the selection operation of the NSGA-II algorithm [14] for selecting appropriate pairs of parents and ensuring the diversity of solutions. We apply two tournament rounds for choosing among candidates. Simply put, an individual has to be fitter than at least two others. CDOXplorer produces two children from two parents via executing a custom crossover (S2) and mutation operator (S3) that are detailed in the Sections V-C and V-D, respectively.

As mentioned before, CDOSim is used for evaluating the individuals in the fourth step S4. In this step, the values of the objectives that have to be optimized are obtained by simulating CDOs. Because the simulations are very time-consuming, we had to limit the population size and configured the basic parameters of CDOXplorer as follows. Our populations contain 50 individuals (population size α). 25 individuals are selected from each generation for reproduction (number of

parents μ) and each generation spawns 50 children (number of children λ). Furthermore, CDOXplorer produces 60 generations per default. Genetic algorithms also use biological analogies for representing the individuals of a population. Their basic elements are specified by *genes*. Considering the classes in Figure 2, the ID of a service represents a single gene, for instance. All genes together constitute the so-called *genome* that contains the complete genetic information of all possible CDOs. Genes can be grouped in larger structures that are called *chromosomes*. Figure 3 illustrates the basic chromosomes and genes that are processed by CDOXplorer. The chromosomes correspond to the class structure of Figure 2 and map to one or more genes that together form a gene sequence. Such a single gene sequence encodes a specific CDO and is called a *genotype*. The *Node Configuration* chromosome constitutes a container for further chromosomes that correspond to classes shown in Figure 2. As there can exist one or more node configurations each having zero or one pair of a *Grow Rule* and *Shrink Rule* chromosome, the genotypes exhibit variable lengths. The *crossover point*s in Figure 3 are detailed in Section V-C. The abbreviations and range of values that are used for the single genes are listed in Table I. These genes correspond to the attributes of classes from Figure 2. We limited their values to a narrow range and discrete spaces for avoiding a further growth of the search-space.

Figure 4 presents three examples of CDOs that are encoded as genotypes. `CDO1` contains two node configurations from which only the second exhibits assigned grow and shrink rules, as can be seen by taking into account the general structure of genotypes in Figure 3. Hence, the cloud environment Amazon EC2 is encoded by the number 7 in this example (gene `CE`) and the first node configuration comprises only the genes 2-4. The deployed service 0 can be identified by the first gene `SE`. The second node configuration that includes a reconfiguration rule is represented by the genes 5 (`SE`) - 20 (`S5`). The further example `CDO2` also includes two node configurations and uses Amazon EC2, whereas `CDO3` shows a genotype using the cloud environment Microsoft Windows Azure and only one node configuration. We will refer to these genotypes later.

B. Problem Statement

After introducing the basic structure of our genetic algorithm, we can describe the multi-objective optimization problem that is tackled by CDOXplorer as follows. Let Φ be the set of all *feasible* CDOs (feasibility is explained below) for a given set of IaaS-based cloud environments and an architectural model, status-quo deployment model, and workload profile of a software system. The goal is to find a CDO $x \in \Phi$ that minimizes the values of the three objective functions $costs(x)$, $rt(x)$, and $sla(x)$ that determine the costs, average response

TABLE I: Design of the used genes.

Gene	Range	Description	Chromosome
CE	$\mathbb{N}$	Cloud environment id	Cloud Env.
SE	$\mathbb{N}$	Service id	Service Comp.
IT	$\mathbb{N}$	VM Instance type id	Initial Start C.
NI	$\mathbb{N}$	Nr. of VM instances to start initially	Initial Start C.
GA	0,1	Grow action; 0: scale-up, 1: scale-out	Grow Rule
G1	$\mathbb{N}$	Minimum nr. of VM instances	Grow Rule
G2	1.1-3.0	MIPIPS multiple in steps of 0.1	Grow Rule
G3	0,1	Condition scope; 0: single VM, 1: all VMs	Grow Rule
G4	0.05-1.0	Condition median utilization in steps of 0.05	Grow Rule
G5	5-60	Condition time period in steps of 5 minutes	Grow Rule
SA	0,1	Shrink action; 0: scale-down, 1: scale-in	Shrink Rule
S1	$\mathbb{N}$	Minimum nr. of VM instances	Shrink Rule
S2	0.1-0.9	MIPIPS multiple in steps of 0.1	Shrink Rule
S3	0,1	Condition scope; 0: single VM, 1: all VMs	Shrink Rule
S4	0.0-0.95	Condition median utilization in steps of 0.05	Shrink Rule
S5	5-60	Condition time period in steps of 5 minutes	Shrink Rule

time, and number of SLA violations of x, respectively. The costs refer to the total amount of monetary units owed to a cloud provider because of utilizing provided services. The response times refer to the average response times of the methods that are included in a workload profile. Lastly, the SLA violations indicate the number of method calls with response times that exceed a given threshold. We denote the set of all node configurations in x as N and the set of all services in x as S. x is feasible if it (1) complies to the structure of CDOs (see Figure 3), (2) complies to the value ranges defined in Table I, and (3) complies to the constraints that are described below. We will use the following notation where y denotes a gene or chromosome and z denotes a chromosome or CDO. Furthermore, gr names a grow rule and sr a shrink rule.

$y \prec z$: y is contained in z
$\triangle(y,z)$: Number of y in z
T_c: Set of VM instance types of cloud environment c

x has to comply with the following constraints (1) - (7).

$$\forall s \in S, x \in \Phi : \triangle(s,x) \geq 1 \tag{1}$$
$$\forall s \in S, n \in N : \triangle(s,n) \leq 1 \tag{2}$$
$$\forall n \in N : \exists s \in S, s \prec n \tag{3}$$
$$\forall IT \prec x, CE \prec x, x \in \Phi : IT \in T_{CE} \tag{4}$$
$$\forall n \in N : \triangle(gr,n) = \triangle(sr,n) \leq 1 \tag{5}$$
$$\forall gr \prec n, sr \prec n, n \in N : gr \succ GA = SA \prec sr \tag{6}$$
$$\forall gr \prec n, sr \prec n, n \in N : gr \succ G4 > S4 \prec sr \tag{7}$$

Constraint (1) describes that each service has to be present at least once in some node configuration of an individual. Furthermore, we do not allow duplicated services in a single node configuration (constraint (2)). The constraint (3) states that at least one service has to be present in each node configuration. A specific VM instance type (gene IT) also has to conform with a stated cloud environment (gene CE, see constraint (4)). Thus, VM instance types of Amazon EC2 cannot be used in conjunction with Microsoft Windows Azure, for instance. Constraint (5) phrases the following limitation: If a grow rule exists in a node configuration, a shrink rule also has to be present in this node configuration and vice versa. Considering grow rules and shrink rules, the grow actions and shrink actions have to match (constraint (6)), i.e., a scale-out rule has to be accompanied by a scale-in rule and a scale-up
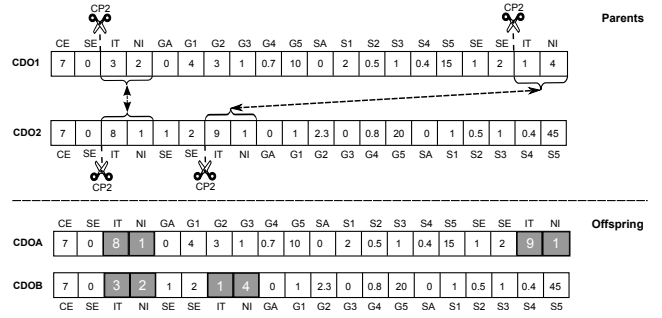

Fig. 5: Crossover operator example.

rule has to be associated with a scale-down rule. The CPU utilization thresholds of grow rules and shrink rules indicate trigger points when to start or shutdown a VM instance. Here, the CPU utilization threshold of a grow rule has to exceed the CPU utilization threshold of a shrink rule (constraint (7)).

C. Crossover Operator

The reproduction procedure involves the application of the crossover operator for producing two children from two parent individuals by mixing their genetic information. This technique follows the biological analogy for passing properties of the parents to their offspring. As the simulation of CDOs is very expensive, we designed the crossover operator to produce only feasible individuals (see Section V-B). Hence, we restricted the mixing to dedicated positions in the genotype called *crossover points*. We defined the four crossover points `CP1-CP4` that are shown in Figure 3. They are selected by chance, are aligned to the chromosomes' boundaries, and specify corresponding gene sequences that can be swapped.

Figure 5 shows an example for applying the crossover operator with the help of the previously introduced CDOs of Figure 4. The example considers `CDO1` and `CDO2`. Both CDOs include two node configurations. As `CP2` was selected, both initial start configurations of `CDO1` and `CDO2` are swapped.

D. Mutation Operator

After two parent individuals have initially produced two children with the help of the crossover operator, the mutation operator is applied to each child. In general, genetic algorithms randomly mutate single genes or gene sequences. This imitates sudden leaps and modifications to the global gene pool that occasionally appear during the evolution process. Considering the influence on the overall optimization procedure, the mutation operator fosters retaining the diversity of the individuals and helps to avoid convergence to a local optimum. Just as the crossover operator, the mutation operator is aligned to the chromosome boundaries that can be seen in Figure 3. As a mutation also has to maintain the inner structure of a chromosome, we divided the mutation operator in five sub operators that are described in the following.

M-CE Mutates the gene CE, i.e., a different cloud environment is used. Here, the IT gene of each node configuration has to be modified as well, as the formerly used VM instance types are not available for the new cloud environment.

M-NN Mutates the number of node configurations and relocates the services. When a node configuration is added,

a service (gene SE) is moved from another node configuration to the new one. When a node configuration is removed, all services are relocated to other node configurations.

M-IS Mutates the initial start configuration of a single node configuration, i.e., another VM instance type (gene IT) is selected or the number of VMs that are initially started for this node configuration (gene NI) is increased or decreased.

M-SC Mutates the service composition of a single node configuration. A service (gene SE) can be added or removed.

M-RR Mutates a reconfiguration rule, i.e., at least one of the genes GA, G1-G5, SA, S1-S5 is modified. When altering a grow rule, changes may also be necessary for the shrink rule to satisfy the constraints and vice versa.

VI. EXPERIMENTAL EVALUATION

To evaluate our simulation-based genetic algorithm CDO-Xplorer, we implemented it as a component of our open source tool CloudMIG Xpress. We let CDOXplorer create CDO candidates on the basis of three well-known cloud environments and investigated the following research questions:

RQ1: Feasibility—Does CDOXplorer reliably provide well-suited results?

RQ2: Competitiveness—Are CDOXplorer's results at least on a par with those from other state-of-the-art search methods?

RQ3: Scalability—Is CDOXplorer applicable for both single and multi cloud environment scenarios?

CDOXplorer utilizes our simulator CDOSim. Please note that the following evaluation covers only CDOXplorer. For an evaluation of CDOSim that demonstrates its applicability and precision, we refer to our previous work [5, 15]. We begin with explaining the methodology that we used to tackle the research questions in Section VI-A. The experimental setting is then described in Section VI-B. The research questions RQ1-RQ3 are detailed in Sections VI-C to VI-E, respectively, before the threats to validity are discussed in Section VI-F.

A. Methodology

We used the Opt4J framework for meta-heuristic optimization [16] as a basis for our implementation of CDOXplorer and also applied the NSGA-II algorithm for the selection of parent individuals (see Section V-A). As NSGA-II utilizes so-called fronts of non-dominated individuals [14], a pareto optimum is also called *pareto optimal front* in this context. There exist several standard performance metrics that basically measure the level of approximation towards the pareto optimal front. As all existing performance metrics exhibit different strengths and weaknesses and therefore are often used in combination [17], we decided against using a single metric. Instead, we applied the two popular metrics *inverted generational distance* (IGD) [18] and *hypervolume indicator* (HV) [19] that are illustrated in Figure 6. Usually, the metrics specify the approximation towards the true pareto optimal front (PF_{true}) if this is known. In our case, PF_{true} can only be obtained by simulating all feasible CDOs. As this is not possible in a reasonable amount of time, we instead use the
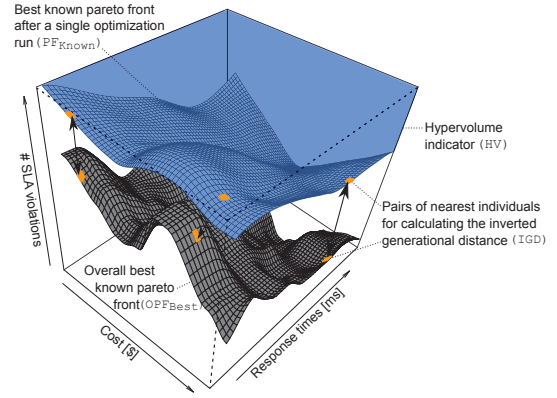

Fig. 6: The metrics *inverted generational distance* and *hypervolume indicator* (colored blue) are used for evaluation.

best pareto optimal front, that is formed by the combination of the overall simulation results, as a proxy and term it OPF_{Best}. Another pareto optimal front depicted in Figure 6 is PF_{Known}. It is the result of each execution of CDOXplorer, whereas each execution includes 3,000 CDOSim simulations. The cube's axes in Figure 6 indicate the three objectives that are optimized by our algorithm: costs, response times, and SLA violations. Starting from the individuals in OPF_{Best}, the nearest individuals in PF_{Known}, in terms of the Euclidean distance, are used for calculating the IGD metric. This metric directly measures the distance from OPF_{Best} to PF_{Known}. Hence, smaller values indicate a better approximation.

In contrast, the values produced by the HV metric become bigger as the approximations become better. HV measures the volume in hyper-dimensional space that is covered by PF_{Known} regarding a reference point $r \in \mathbb{R}^d$ in d-dimensional space. As we consider three objectives and configure HV accordingly, the mentioned volume corresponds to the volume in three-dimensional space that is colored blue in the illustration of Figure 6. HV therefore implicitly rates the distance from OPF_{Best} to PF_{Known} as is also done by IGD. Additionally, both metrics assess the spread of PF_{Known}, that constitutes another important quality characteristic. When covering larger parts of the search space, there exist less unexplored areas that could potentially contain better solutions.

To judge the measurement results of these metrics for CDOXplorer, we used several state-of-the-art search and optimization techniques for comparison. However, as the objective values cannot be obtained by solving functions analytically, some popular classes of approaches, such as gradient-based optimization methods, cannot be used. Though, direct search methods are suited for simulation-based optimization [20]. Therefore, we use the simple yet effective stochastic algorithms *simple random sampling (SI-RS)* and *systematic random sampling (SY-RS)* [21] as two out of the three algorithms used for comparison with CDOXplorer. The algorithm SI-RS creates 3,000 CDO individuals by chance and serves as a baseline algorithm. SY-RS also produces 3,000 individuals by chance, but works in a different way. It can iterate over all f feasible CDO candidates and randomly selects the c-th CDO from the $[0, floor(\frac{f}{3,000})]$ interval at the beginning of

Fig. 7: Simulated annealing reuses mutation sub operators.

a CDOXplorer run, where $floor(x)$ rounds $x \in \mathbb{R}_+$ down to the next natural number. The next CDO is then given by the $(c + floor(\frac{f}{3,000}))$-th candidate and so forth.

The third algorithm used for comparison purposes is, just as a genetic algorithm, a nature-inspired meta-heuristic termed *simulated annealing (SI-AN)* [22]. Basically, it mimics the temperature cooling process of materials. To emulate such a cooling process in our problem context, we reuse the mutation sub operators introduced in Section V-D as illustrated in Figure 7. The temperature is adapted according to five phases. The phases utilize mutation sub operators with specific probabilities to reduce disruptive modifications over time.

CDOXplorer uses 60 generations with populations of 50 individuals. Hence, CDOXplorer applies 3,000 simulations in a single run. The runs for CDOXplorer, SI-RS, SY-RS, and SI-AN were each repeated 40 times for single as well as for multi cloud scenarios. Thus, we conducted 320 optimization runs with 960,000 simulations in total. We therefore obtained 480,000 simulations for the single and multi cloud scenario that were used to approximate two pareto optimal fronts that were set to OPF_{Best} for evaluating the respective scenarios. The first of these pareto optimal fronts is shown in Figure 8.

B. Experimental Setting

For incorporating public clouds, we measured the MIPIPS values (see Section III) for 12 VM instance types of Amazon EC2, and five of Microsoft Windows Azure, by using locations in Europe for both. Five VM instance types of our private Eucalyptus cloud were also benchmarked. For analyzing single and multi cloud scenarios, we used the resulting cloud profiles in the two corresponding scenarios SC_S and SC_M as follows.

SC$_S$: Amazon EC2

SC$_M$: Amazon EC2, Microsoft Windows Azure, Eucalyptus

We extracted a KDM model from the open source ERP system Apache OfBiz 10.04.[2] Then, we deployed this system on a machine of our local cluster and described the deployment in a status-quo deployment model. Customers that browse the webstore and put articles in their shopping carts were emulated by producing workload according to a typical day/night usage pattern. More customers visited the webstore in the evening instead in the morning hours and the traffic largely reduced at night. The measured response times and the MIPIPS value of our hardware were then used in a workload profile for generating CDOs and driving the CDOSim simulations. The SLA violation threshold was set to 2s. For Eucalyptus, we defined a synthetic cost model where the prices for VMs follow the capabilities of our VM instance types.

[2] http://ofbiz.apache.org/

Fig. 8: One of two best known pareto fronts OPF_{Best}.

C. RQ1: Feasibility

This research question evaluates the feasibility and applicability of CDOXplorer. As a basic requirement, CDOXplorer has to reliably provide well-suited results. The evaluation of this criteria is of particular importance because of the following two reasons. (1) The simulations used for our simulation-based algorithm are computationally expensive. Hence, we strictly limited the number of generations and population size. This could affect CDOXplorer's capability for producing well-suited approximations of pareto optimal fronts. (2) The non-determinism used in CDOXplorer, for example, regarding the selection of crossover points, could possibly lead to considerable variations among optimization runs. We approach RQ1 by computing two further metrics M1 and M2.

M1 analyzes the quality of the results produced by CDO-Xplorer. We were interested in the degree the hypervolumes of OPF_{Best} can be approximated by CDOXplorer for SC_S and SC_M. For SC_S and SC_M, the HV of OPF_{Best} was 0.462 and 0.573, respectively, whereas CDOXplorer achieved 0.448 and 0.565. Thus, the actual quality of the found pareto optima were sufficiently well-suited for our needs, as these results turn into 96.96% and 98.56% approximation of OPF_{Best} for the single and multi cloud scenario SC_S and SC_M, respectively.

The second metric M2 calculates the coefficient of variation (CV) that gives information about the relative dispersion regarding a sample's mean value μ. As our performance metrics IGD and HV deliver results in artificial and incomparable units, we convert the standard deviation σ for SC_S and SC_M in combination with IGD and HV to relative and therefore comparable values. These are denoted as CV_S for our single and CV_M for our multi cloud scenario. In general, CV is computed by $CV = \frac{\sigma}{\mu}$. The IGD results for SC_S and SC_M vary in a band of 7.77% and 15.84% around μ, and of 0.46% and 0.32% around μ for HV, respectively. Hence, IGD results in the single cloud scenario are up to 3.85% lower or higher than μ, for instance. Thus, the results indicate that CDOXplorer can reliably find well-suited solutions. However, the value for IGD increases for a higher number of cloud profiles. In our future work, we will investigate if this constitutes an actual trend.

D. RQ2: Competitiveness

This research question addresses the competitiveness with other state-of-the-art search approaches by comparing CDO-Xplorer with SI-RS, SY-RS, and SI-AN. Table II lists the results for the metrics IGD and HV for the single cloud scenario SC_S. Table III shows these results for SC_M. The tables list the

TABLE II: Search Over A Single Cloud Profile (SC_S).

Metric		Search Method			
		CDOXplorer	SI-RS	SY-RS	SI-AN
I.G. Distance	Mean	**2.70E-02**	3.67E-02	4.11E-02	3.28E-02
	SD	2.10E-03	2.13E-03	3.61E-03	2.85E-03
	Median	**2.72E-02**	3.65E-02	4.21E-02	3.20E-02
	Min (best)	2.16E-02	3.34E-02	3.40E-02	2.76E-02
	Max (worst)	3.03E-02	4.07E-02	4.83E-02	3.95E-02
Hypervolume	Mean	**4.48E-01**	4.41E-01	4.41E-01	4.44E-01
	SD	2.08E-03	1.96E-03	2.89E-03	2.09E-03
	Median	**4.48E-01**	4.40E-01	4.41E-01	4.44E-01
	Min (worst)	4.44E-01	4.36E-01	4.35E-01	4.40E-01
	Max (best)	4.54E-01	4.46E-01	4.46E-01	4.48E-01

TABLE III: Search Over Three Cloud Profiles (SC_M).

Metric		Search Method			
		CDOXplorer	SI-RS	SY-RS	SI-AN
I.G. Distance	Mean	**3.08E-02**	7.18E-02	8.03E-02	3.37E-02
	SD	4.88E-03	2.78E-03	4.41E-03	4.50E-03
	Median	**3.12E-02**	7.17E-02	7.90E-02	3.38E-02
	Min (best)	2.13E-02	6.58E-02	7.23E-02	2.52E-02
	Max (worst)	4.16E-02	7.76E-02	8.88E-02	4.67E-02
Hypervolume	Mean	**5.65E-01**	5.20E-01	5.18E-01	5.63E-01
	SD	1.82E-03	1.95E-03	2.43E-03	1.68E-03
	Median	**5.65E-01**	5.20E-01	5.17E-01	5.63E-01
	Min (worst)	5.61E-01	5.16E-01	5.13E-01	5.61E-01
	Max (best)	5.70E-01	5.25E-01	5.25E-01	5.68E-01

Fig. 9: CDOXplorer advantage relative to other approaches.

mean, standard deviation, median, min, and max values of 40 complete, repeated optimization runs for each combination of metric, scenario, and search method. Bigger values are better for HV but worse for IGD. All best mean and median values are set in bold. As can be seen, CDOXplorer outperforms all other search methods in SC_S and also in SC_M.

We use the Mann-Whitney non-parametric test to evaluate statistical significance. The null hypothesis H_0 states that the results from CDOXplorer cannot be distinguished from those of SI-RS, SY-RS, and SI-AN. Using the Mann-Whitney test and a Bonferroni correction ($\alpha_1 = 0.016$) for multiple comparisons, H_0 is rejected with significance level (α) 0.05 for all combinations of metrics with SC_S and SC_M. Thus, we can quantify the degree CDOXplorer performs better and compare the medians of SI-RS, SY-RS, and SI-AN to those of CDOXplorer. Figure 9 shows the corresponding fractions with regard to the medians, e.g., for the SC_M scenario, the median for IGD is over 60% lower than that of SY-RS.

E. RQ3: Scalability

This research question analyzes CDOXplorer's scalability. That means, if CDOXplorer can retain its performance when more cloud profiles are considered for providing CDO candidates or if a potential performance degradation is still acceptable. This is of particular interest as the number of generations and the population size remain stable but the search-space size grows linearly with each new cloud profile.

As described in Section VI-C, the value for the coefficient of variation (with IGD) and the value for the IGD metric itself grows when CDOXplorer processes SC_M. However, the HV metric value and the coefficient of variation even become better, when more cloud profiles are used. Furthermore, it

is useful to compare these observations with those from the other search methods. CDOXplorer provides in all cases the best results. Considering SI-RS and SY-RS, their values for the IGD metric grow even more and nearly doubled when using SC_M. SI-AN also suffers from a deterioration when transitioning from the SC_S to the SC_M scenario. Though, the IGD values increase less. The values for HV show a similar development. CDOXplorer is better than all other search methods and the scalability is better compared with SI-RS and SY-RS, but SI-AN exhibits a slightly higher improvement.

In summary, we conclude that (1) CDOXplorer scales better than SI-RS and SY-RS, but (2) slightly worse than SI-AN. Although CDOXplorer performs better in all cases in terms of the absolute median values, we will conduct additional experiments in our future work to further investigate scalability.

F. Threats to Validity

There are several issues that form a threat to validity. First, we only consider three cloud profiles, as their construction is not trivial and the VM instance types of additional cloud providers need to be benchmarked. This involves potential expenses. However, as indicated by the results of Section VI-E, scalability is a worthwhile area for further analyses.

Restrictions were also made concerning the workload profile and the studied enterprise software. In each case, just one sample was used. This is due to the fact that further optimization runs imply even more time-consuming simulations. Therefore, we strived after using representative instances. Day/night usage patterns with higher and lower demand are frequently found for enterprise systems. Furthermore, Apache OfBiz is very popular and widespread. However, we plan to address the stated limitations in our future work.

There could also exist other optimization methods that provide better results than SI-RS, SY-RS, and SI-AN we used for evaluation. Though, tailoring optimization methods for our context is time-consuming and not straight forward. Quite similarly, there could exist better ways to tailor SI-AN instead of reusing the mutation sub operators. Especially in the light that SI-AN partially comes near to CDOXplorer's results.

Further threats to validity arise from the synthetic cost model for our private Eucalyptus cloud and that SLAs are

usually defined in terms of percentile ranges when considering response times. However, aligning the VM prices to the capabilities of the VM instance types is omnipresent with respect to public cloud environments. Altering the absolute threshold into percentiles for defining the SLA objective can be easily done and will be addressed in the future work.

VII. Related Work

This section discusses related work concerning the optimization of deployment architectures and reconfiguration rules. Here, solely approaches that tackle related problems from a user perspective are considered. Cloud users want to deploy software to the cloud under given constraints. Especially, a cloud environment's internal structure is transparent to cloud users. As there exists a large body of work in this regard, we limit the description to selected approaches.

Deployment optimization in non-cloud scenarios. For improving the deployment architecture of distributed systems regarding arbitrary QoS properties, Malek et al. [23] propose an extensible framework and visual modeling and analysis environment. The framework incorporates continuous system monitoring and changing the deployment architecture at run-time by actually executing redeployment operations. A further mode allows for offline simulation. Arbitrary deployment constraints and QoS properties can be formally defined by manually specifying utility functions. The approach provides four predefined deployment improvement algorithms, among those is a genetic algorithm. The utility functions are used as fitness functions, whereas CDOXplorer uses simulation runs to obtain fitness values. Furthermore, Malek et al. [23] do not consider the integration of reverse-engineered code models and monitored workload, as is supported by our approach.

Martens et al. [24] focus on component-based systems and at finding optimized software and deployment architectures concerning performance, reliability, and costs. Similar to our approach, a genetic algorithm is used and simulations are, partially, employed for assessing candidate solutions. Their software PerOpteryx provides tool support and explores four degrees of freedom, e.g., processor speeds. No dynamic resource scaling is supported, but further degrees of freedom can be added. In contrast, CDOXplorer currently explores 17 fixed degrees of freedom (the number of node configurations and 16 genes) and optimizes runtime reconfiguration rules.

A hybrid approach that considers the consolidation of multi-tier applications to virtual machines is presented by Jung et al. [25]. Applications are modeled and optimized offline to generate well-suited system configurations. These are transformed to adaptation policies that can be consumed online by rule engines. Compared with CDOXplorer, the approach by Jung et al. only allows to define a single resource type and homogeneous resource instances instead of considering distinct VM instance types that can also be used in parallel.

Zhang et al. [26] propose a QoS-aware approach for finding the optimal number of machines for deploying services in the context of service oriented architectures (SOAs). A greedy algorithm is used that maximizes the throughput. Unlike in our approach, dynamic resource scaling is not supported.

Non-evolutionary cloud deployment optimization. As with our approach CloudMIG, Wu et al. [27] consider resource usage optimization for SaaS providers that build upon leased VMs. Wu et al. assume that a maximum service utilization is defined per customer in SLA agreements. SLA violations and infrastructure costs are minimized by two custom-made algorithms that are evaluated with the tool CloudSim. Our simulator CDOSim also builds on CloudSim. In contrast to CDOXplorer, Wu et al. do not consider distributed applications, runtime reconfiguration, and arbitrary workload.

Trummer et al. [28] contribute an algorithm for computing an application's cost-optimal deployment architecture for IaaS-based clouds. This single-objective optimization is described as a constraint optimization problem and is tackled with an existing constraint solver. As opposed to this, we use multi-objective optimization and support dynamic resource scaling. Nevertheless, the application templates used by Trummer et al. serve the same purpose as our status-quo deployment models.

San Aniceto et al. [29] also use a custom-made single-objective optimization algorithm for reducing the costs of leased VMs. Besides on-demand instances—that can be started and stopped at any time—it also considers reserved instances. Those are frequently offered for a single payment per time period. In turn, a discount is given in the hour rates. The algorithm aims at finding the best suited combination of on-demand and reserved instances based on historical workload data. However, it does not support dynamic reconfiguration.

Dynamic scaling is, in contrast, supported by the approach of Mao et al. [30]. Here, integer programming problems are solved for reducing costs or maximizing the performance of scientific computing jobs. Scaling policies are derived at runtime by learning from previous job executions. Compared with our approach, we target enterprise software and use simulations for assessing different CDO candidates before actually deploying an application to a specific cloud environment.

Evolutionary cloud deployment optimization. Most related work that employs evolutionary optimization techniques in cloud deployment scenarios takes a cloud provider perspective or requires corresponding knowledge regarding the internal structure of a cloud environment (e.g., [31, 32]). In contrast, Wada et al. [33] follow, in common with Cloud-MIG, a cloud user perspective and contribute the genetic algorithm E^3-R that explores deployment configurations for optimizing services' QoS attributes. E^3-R can also reduce redundant QoS objectives and estimate the performance of configurations using queueing theory and historic mean arrival rates. However, E^3-R does not support varying workload and dynamic resource scaling, and it regards services as black boxes, whereas our approach simulates reverse-engineered and transformed architectural models and monitored workload.

A particle swarm optimization-based heuristic is introduced by Pandey et al. [34]. The single-objective optimization algorithm maps scientific tasks to cloud resources for minimizing costs. In comparison, we target optimal CDOs for enterprise software, formulate the search as a multi-objective optimization problem, and facilitate dynamic resource scaling.

VIII. Conclusion

We presented our simulation-based genetic algorithm CDO-Xplorer that optimizes cloud deployment options (CDOs) for supporting the migration of software to the cloud. A CDO comprises a deployment architecture and runtime reconfiguration rules that can start further VM instances when the overall utilization exceeds a given threshold, for instance. Finding near-optimal solutions is approached by using techniques of the search-based software engineering field. CDOs become optimized in terms of response times, costs, and number of SLA violations. For example, the best-suited cloud provider and VM instance types for an existing software system and particular usage patterns have to be found. CDOXplorer uses our tool CDOSim to simulate CDOs. A simulation result represents a fitness value of our genetic algorithm.

The popular public clouds Amazon EC2 and Microsoft Windows Azure, as well as our private Eucalyptus cloud, were used for extensive experiments. We compared CDOXplorer with three state-of-the-art search and optimization methods. Our evaluation showed that CDOXplorer can find CDOs that are up to 60% better than those of the other approaches. Our experiment code and data and an implementation of CDOXplorer are available online as open source software.

References

[1] M. Armbrust, A. Fox, R. Griffith, A. D. Joseph, R. Katz, A. Konwinski, G. Lee, D. Patterson, A. Rabkin, I. Stoica, and M. Zaharia, "A view of cloud computing," *Commun. ACM*, vol. 53, no. 4, pp. 50–58, Apr. 2010.

[2] S. Marston, Z. Li, S. Bandyopadhyay, J. Zhang, and A. Ghalsasi, "Cloud computing – the business perspective," *Decision Support Systems*, vol. 51, no. 1, pp. 176–189, 2011.

[3] J. Grundy, G. Kaefer, J. Keong, and A. Liu, "Guest Editors' Introduction: Software Engineering for the Cloud," *IEEE Software*, vol. 29, 2012.

[4] G. Canfora, M. Di Penta, R. Esposito, and M. L. Villani, "An approach for QoS-aware service composition based on genetic algorithms," in *Proc. of the Conf. on Genetic and Evol. Computation*. ACM, 2005.

[5] F. Fittkau, S. Frey, and W. Hasselbring, "CDOSim: Simulating Cloud Deployment Options for Software Migration Support," in *Proceedings of the 6th International Workshop on the Maintenance and Evolution of Service-Oriented and Cloud-Based Systems (MESOCA 2012)*. IEEE Computer Society, 2012, pp. 37–46.

[6] S. Frey, W. Hasselbring, and B. Schnoor, "Automatic conformance checking for migrating software systems to cloud infrastructures and platforms," *Journal of Software: Evolution and Process*, doi: 10.1002/smr.582, 2012.

[7] S. Frey and W. Hasselbring, "The CloudMIG Approach: Model-Based Migration of Software Systems to Cloud-Optimized Applications," *Int'l Journal on Advances in Software*, vol. 4, no. 3 and 4, 2011.

[8] M. Harman, "Software Engineering Meets Evolutionary Computation," *Computer*, vol. 44, no. 10, pp. 31–39, 2011.

[9] P. Mell and T. Grance, "The NIST Definition of Cloud Computing," Sept. 2011, NIST Special Publication 800-145.

[10] A. Law and M. McComas, "Simulation-based optimization," in *Proc. of the Winter Simulation Conference, 2000*, vol. 1, 2000, pp. 46–49.

[11] H. Bruneliere, J. Cabot, F. Jouault, and F. Madiot, "MoDisco: A Generic And Extensible Framework For Model Driven Reverse Engineering," in *Proc. of the IEEE/ACM Int'l Conference on Automated Software Engineering*, ser. ASE '10, 2010.

[12] R. Pérez-Castillo, I. G.-R. de Guzmán, and M. Piattini, "Knowledge Discovery Metamodel-ISO/IEC 19506: A standard to modernize legacy systems," *Computer Standards & Interfaces*, vol. 33, no. 6, 2011.

[13] A. van Hoorn, J. Waller, and W. Hasselbring, "Kieker: A framework for application performance monitoring and dynamic software analysis," in *Proc. of the 3rd ACM/SPEC Int'l Conference on Performance Engineering (ICPE 2012)*. ACM, Apr. 2012, pp. 247–248.

[14] K. Deb, A. Pratap, S. Agarwal, and T. Meyarivan, "A fast and elitist multiobjective genetic algorithm: NSGA-II," *IEEE Transactions on Evolutionary Computation*, vol. 6, no. 2, pp. 182 –197, Apr. 2002.

[15] F. Fittkau, S. Frey, and W. Hasselbring, "Cloud User-Centric Enhancements of the Simulator CloudSim to Improve Cloud Deployment Option Analysis," in *Proceedings of the European Conference on Service-Oriented and Cloud Computing (ESOCC)*, 2012, (to appear).

[16] M. Lukasiewycz, M. Glaß, F. Reimann, and J. Teich, "Opt4J: A Modular Framework for Meta-Heuristic Optimization," in *Proc. of the 13th Annual Conf. on Genetic and Evolutionary Computation*. ACM, 2011.

[17] E. Zitzler, L. Thiele, M. Laumanns, C. Fonseca, and V. da Fonseca, "Performance assessment of multiobjective optimizers: an analysis and review," *IEEE Transactions on Evolutionary Computation*, vol. 7, no. 2, pp. 117–132, Apr. 2003.

[18] Q. Zhang, A. Zhou, and Y. Jin, "RM-MEDA: A Regularity Model-Based Multiobjective Estimation of Distribution Algorithm," *IEEE Transactions on Evolutionary Computation*, vol. 12, no. 1, pp. 41–63, Feb. 2008.

[19] E. Zitzler and L. Thiele, "Multiobjective optimization using evolutionary algorithms - a comparative case study," in *Parallel Problem Solving from Nature*, ser. Lecture Notes in Comp. Science. Springer, 1998, vol. 1498.

[20] T. G. Kolda, R. M. Lewis, and V. Torczon, "Optimization by Direct Search: New Perspectives on Some Classical and Modern Methods," *SIAM Review*, vol. 45, pp. 385–482, 2003.

[21] F. Schoen, "Stochastic techniques for global optimization: A survey of recent advances," *Journal of Global Optimization*, vol. 1, 1991.

[22] C. Blum and A. Roli, "Metaheuristics in combinatorial optimization: Overview and conceptual comparison," *ACM Comput. Surv.*, vol. 35, no. 3, pp. 268–308, Sep. 2003.

[23] S. Malek, N. Medvidovic, and M. Mikic-Rakic, "An Extensible Framework for Improving a Distributed Software System's Deployment Architecture," *IEEE Transactions on Software Engineering*, vol. 38, no. 1, pp. 73–100, jan.-feb. 2012.

[24] A. Martens, H. Koziolek, S. Becker, and R. Reussner, "Automatically Improve Software Architecture Models for Performance, Reliability, and Cost Using Evolutionary Algorithms," in *Proceedings of the First Joint WOSP/SIPEW International Conference on Performance Engineering*. ACM, 2010, pp. 105–116.

[25] G. Jung, K. Joshi, M. Hiltunen, R. Schlichting, and C. Pu, "Generating Adaptation Policies for Multi-tier Applications in Consolidated Server Environments," in *Int'l Conf. on Autonomic Computing, 2008*, june 2008.

[26] C. Zhang, R. Chang, C.-S. Perng, E. So, C. Tang, and T. Tao, "QoS-Aware Optimization of Composite-Service Fulfillment Policy," in *IEEE Int'l Conference on Services Computing, 2007*, july 2007, pp. 11–19.

[27] L. Wu, S. Garg, and R. Buyya, "SLA-Based Resource Allocation for Software as a Service Provider (SaaS) in Cloud Computing Environments," in *11th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGrid), 2011*, May 2011, pp. 195–204.

[28] I. Trummer, F. Leymann, R. Mietzner, and W. Binder, "Cost-optimal outsourcing of applications into the clouds," in *IEEE Second Int'l Conf. on Cloud Computing Technology and Science (CloudCom), 2010*, 2010.

[29] I. San Aniceto, R. Moreno-Vozmediano, R. Montero, and I. Llorente, "Cloud capacity reservation for optimal service deployment," in *Proceedings of the Second International Conference on Cloud Computing, GRIDs, and Virtualization*. IARIA Conference, Sept. 2011, pp. 52–59.

[30] M. Mao, J. Li, and M. Humphrey, "Cloud Auto-scaling with Deadline and Budget Constraints," in *Proceedings of the 11th IEEE/ACM International Conference on Grid Computing (GRID)*, 2010, pp. 41–48.

[31] M. J. Csorba, H. Meling, and P. E. Heegaard, "Ant system for service deployment in private and public clouds," in *Proc. of the 2nd Workshop on Bio-inspired Algorithms for Distributed Systems*. ACM, 2010.

[32] Z. I. M. Yusoh and M. Tang, "Composite SaaS Placement and Resource Optimization in Cloud Computing Using Evolutionary Algorithms," *IEEE Fifth Int'l Conference on Cloud Computing*, pp. 590–597, 2012.

[33] H. Wada, J. Suzuki, Y. Yamano, and K. Oba, "Evolutionary deployment optimization for service-oriented clouds," *Software: Practice and Experience*, vol. 41, no. 5, pp. 469–493, 2011.

[34] S. Pandey, L. Wu, S. M. Guru, and R. Buyya, "A Particle Swarm Optimization-Based Heuristic for Scheduling Workflow Applications in Cloud Computing Environments," in *Proc. 24th IEEE Int'l Advanced Information Networking and Applications (AINA) Conference*, 2010.